AI usage log

Tools. Claude (chat): research, planning, specifications, review of every output, and the written deliverables. GitHub Copilot in Agent mode: most of the code and tests, always working from a specification written in Claude. Gemini: the one AI component that runs inside the app (reading documents and writing the plain-English explanation). The build did not use Claude Code or Claude Desktop. How I worked: specify with numbers to hit, let the tool build in small phases, then check the result against those numbers myself.

Phase	Tool and key prompt (short)	What worked	What I corrected	Time saved (est.)
Plan	Claude: "Review this UX plan against deliverables D1 to D6 and list the gaps"	Found that the research note and AI log had no home and that sanction and disbursement were missing; settled on one orchestrator agent	Cut a persona switcher, coach marks and saved filters to fit the week	~3 h
Research	Claude with web search: stages, documents, FOIR, bureau and RBI rules	Cited sources in one pass	Two sources disagreed on the weekly bureau-reporting date, so the note says "unconfirmed"	~3 h
Data	Copilot, from a Claude spec: 500 synthetic customers and six personas with target FOIR and scores	Deterministic data, fake IDs and watermarked PDFs	I recalculated EMIs by hand; the document-mismatch persona's FOIR is 40.9% on the payslip but 62.0% on bank credits, which made her story honest	~6 h
Agent	Copilot: policy rules as code; Gemini only for extraction and wording	Every decision traceable to a policy clause	The first "recorded" Gemini text was pre-written by the tool, so I relabelled it "pre-written sample". The tool also relaxed one rule (current-job tenure) to produce a decline I wanted for the override scene; I kept it and documented it	~6 h
UI	Screenshots to Claude for review	Spotted invisible input boxes, the verdict buried under process detail, and wrapped step cards; led to a decision-first redesign	Rewrote the brief twice	~4 h
Bugs	Copilot with a root-cause-first prompt	Found and fixed a form crash (state lost between steps)	A review action saved "Reject" as "Approve"; I made the tool reproduce and test it before changing code	~2 h

What worked. Writing the expected numbers into each prompt (for example "FOIR within 2 points"); asking for a DONE / NOT DONE list after each phase; keeping policy

decisions in plain code so the AI cannot change a decision. What did not. Tools reported success without proving it (no live Gemini key, untested browser clicks); I asked for evidence each time. Some figures the tool produced, such as the baseline-versus-demo turnaround, are illustrative and are labelled that way in the app. Key prompt (design rule). "Policy decisions are deterministic code. Gemini only extracts document fields and writes the explanation; its output can never change the decision. The AI never sanctions."